

MQTT

IoT Protocols MQTT

Stefan Forsström

Institution of Information Systems and Technology



Overview

- History and Background
- Structure and Features
- MQTT Messages Details
- Tips for the Lab Project

History

- Message Queue Telemetry Transport (MQTT)
 - Was invented in 1999 by
 - Andy Stanford-Clark (IBM)
 - Arlen Nipper (Arcom, now Cirrus Link)
- The use case was to create a protocol for minimal battery loss and minimal bandwidth, to connect oil pipelines over satellite Internet
- They specified the following goals of the protocol:
 - Simple to implement
 - Provide a Quality of Service Data Delivery
 - Lightweight and Bandwidth Efficient
 - Data Agnostic
 - Continuous Session Awareness

History

- It is an publish-subscribe "lightweight" messaging protocol
 - But it fits the IoT very well
- Highly centralized with a coordinating broker server
 - Consumers subscribe to topics
 - Which producers can publish to
 - And the broker unicasts the data to the subscribers
- Not really “lightweight”, since it is based on TCP
 - Comparable to REST, about the same
 - But has built in publish subscribe features

Structure

TCP/IP Port: 1883

When running over SSL/TLS port: 8883

When running over Websockets port: 8000

Application

MQTT

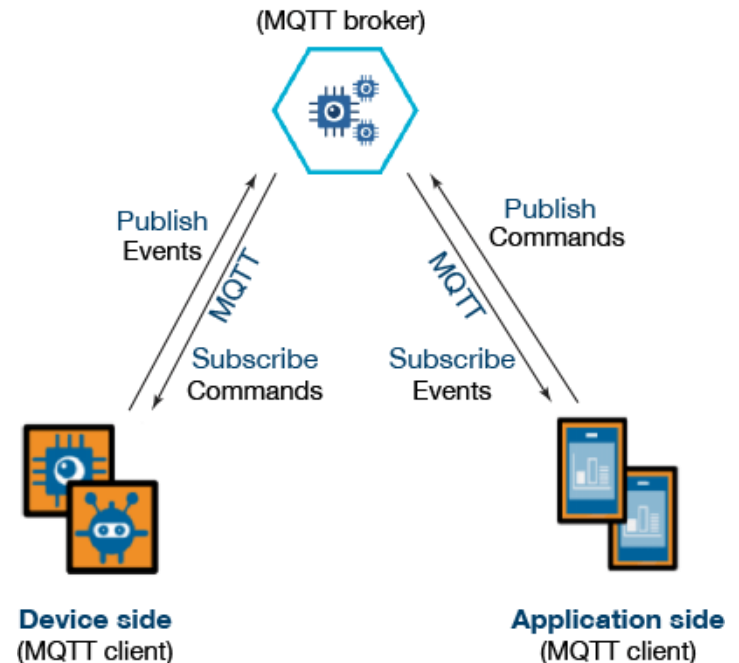
SSL/TLS
optional

TCP

IP

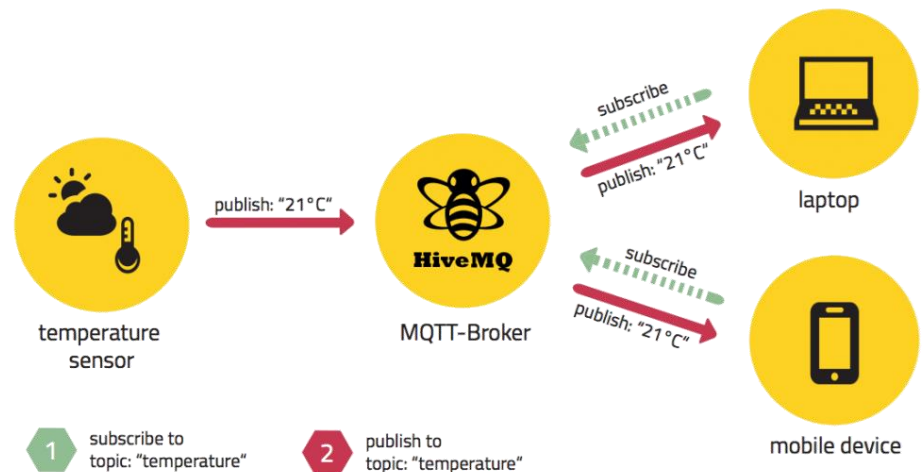
Structure

- MQTT consist of three parts:
 - Broker
 - Subscribers
 - Publishers
- Clients connect to a “Broker”
- Clients subscribe to topics e.g.,
- Clients can publish messages to topics:
 - All clients receive all messages published to topics they subscribe to
- Messages can be anything, Text, Images, etc.



Publish/Subscribe Concept

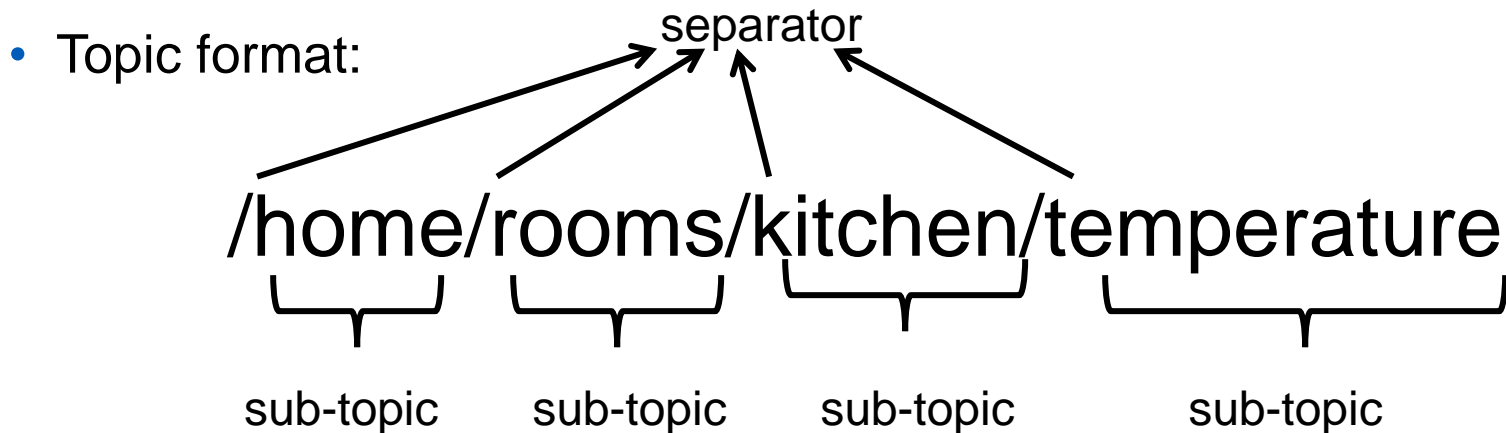
- Decoupled in space and time:
 - The clients does not need to know each others IP address and port and they do not need to be running at the same time



- The broker's IP and port must however be known by all clients

Topics

- Each published data specifies a topic
 - Each subscriber subscribed to that topic will receive it
 - Namespace hierarchy is used for topic filtering



Subscriptions

- Subscription types
 - Durable
 - If the subscriber disconnect messages are buffered at the broker and delivered upon reconnection
 - Non-durable
 - Connection lifetime is the subscription lifetime
 - When the TCP session breaks down, the subscription is closed

Publishing

- It might also be the case that a published message never becomes consumed by any subscriber
 - The clients are unaware of the number of subscribers
- Message retention
 - Retained (a type of “persistent” message)
 - The subscriber upon first connection receives the last good publication (i.e., does not have to wait for new publication)
 - Only the most recent persistent message is stored and distributed
- Last Will and Testament (LWT)
 - A message published upon disconnecting a connection
 - Anybody subscribing to the LWT topic will know when a certain device (that registered a LWT) disconnected


Publishing “QoS” (Reliability)

- 0 – unreliable (aka “at most once”)
 - OK for continuous streams, least overhead (1 message)
 - “Fire and forget”
 - TCP will still provide reliability




Publishing “QoS” (Reliability)

- 1 – delivery “at least once” (duplicates possible)
 - Used for alarms – more overhead (2 messages)
 - Contains message ID (to match with ACKed message)



Publishing “QoS” (Reliability)

- 2 – delivery “exactly once”
 - Utmost reliability is important – most overhead (4 messages) and slowest



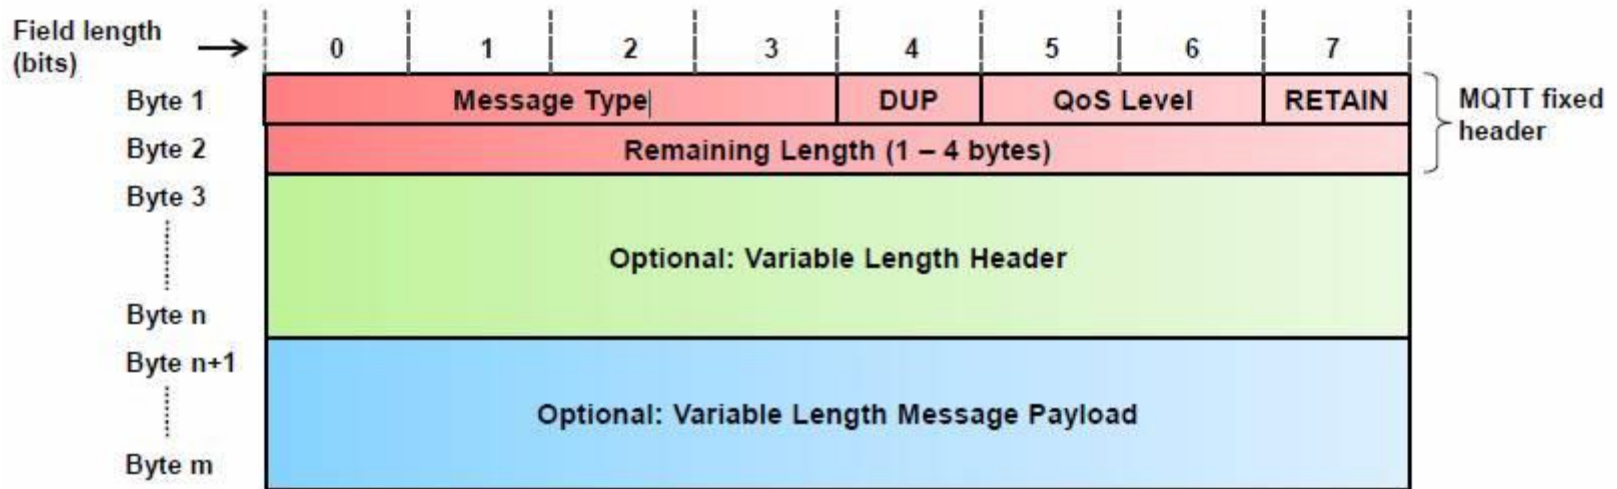
Security?

- All communication is done in clear text
 - Unless SSL/TLS is used
- There is a simple Client ID method used for recognizing users
- But there are also functions for username/password authentication
 - For private accessing the broker etc.
- You can also control which clients are able to subscribe and publish to different topics
 - Using either the ClientID or username/password




MQTT Messages

MQTT Message Format




- Shortest message is two bytes (red fields)



Message Types

Message fixed header field	Description / Values	
Message Type	0: Reserved	8: SUBSCRIBE
	1: CONNECT	9: SUBACK
	2: CONNACK	10: UNSUBSCRIBE
	3: PUBLISH	11: UNSUBACK
	4: PUBACK	12: PINGREQ
	5: PUBREC	13: PINGRESP
	6: PUBREL	14: DISCONNECT
	7: PUBCOMP	15: Reserved
DUP	Duplicate message flag. Indicates to the receiver that this message may have already been received. 1: Client or server (broker) re-delivers a PUBLISH, PUBREL, SUBSCRIBE or UNSUBSCRIBE message (duplicate message).	
QoS Level	Indicates the level of delivery assurance of a PUBLISH message. 0: At-most-once delivery, no guarantees, «Fire and Forget». 1: At-least-once delivery, acknowledged delivery. 2: Exactly-once delivery. Further details see MQTT QoS .	
RETAIN	1: Instructs the server to retain the last received PUBLISH message and deliver it as a first message to new subscriptions. Further details see RETAIN (keep last message) .	
Remaining Length	Indicates the number of remaining bytes in the message, i.e. the length of the (optional) variable length header and (optional) payload. Further details see Remaining length (RL) .	




Connect Message

- Message nr 1
- Protocol name
- Flags
- Keep alive
- Payload
 - can include client ID length and name

Bit	7	6	5	4	3	2	1	0
Byte 1	MQTT Control Packet type (1)				Reserved			
	0	0	0	1	0	0	0	0
Byte 2...	Remaining Length							

	Description	7	6	5	4	3	2	1	0
Protocol Name									
byte 1	Length MSB (0)	0	0	0	0	0	0	0	0
byte 2	Length LSB (4)	0	0	0	0	0	1	0	0
byte 3	'M'	0	1	0	0	1	1	0	1
byte 4	'Q'	0	1	0	1	0	0	0	1
byte 5	'T'	0	1	0	1	0	1	0	0
byte 6	'T'	0	1	0	1	0	1	0	0
byte 7	Level(4)	0	0	0	0	0	1	0	0
Bit	7	6	5	4	3	2	1	0	
	User Name Flag	Password Flag	Will Retain	Will QoS		Will Flag	Clean Session	Reserved	
Byte 8	X	X	X	X	X	X	X	0	
Bit	7	6	5	4	3	2	1	0	
byte 9	Keep Alive MSB								
byte 10	Keep Alive LSB								



Connect Ack Message

- Message nr 2


Bit	7	6	5	4	3	2	1	0
byte 1	MQTT Control Packet Type (2)				Reserved			
	0	0	1	0	0	0	0	0
byte 2	Remaining Length (2)							
	0	0	0	0	0	0	1	0

- SP: Session Present Bit

	Description	7	6	5	4	3	2	1	0
Connect Acknowledge Flags	Reserved								SP ¹
byte 1		0	0	0	0	0	0	0	X
Connect Return code									
byte 2		X	X	X	X	X	X	X	X

- Return code

Value	Return Code Response	Description
0	0x00 Connection Accepted	Connection accepted
1	0x01 Connection Refused, unacceptable protocol version	The Server does not support the level of the MQTT protocol requested by the Client
2	0x02 Connection Refused, identifier rejected	The Client identifier is correct UTF-8 but not allowed by the Server
3	0x03 Connection Refused, Server unavailable	The Network Connection has been made but the MQTT service is unavailable
4	0x04 Connection Refused, bad user name or password	The data in the user name or password is malformed
5	0x05 Connection Refused, not authorized	The Client is not authorized to connect
6-255		Reserved for future use




Publish Message

- Message nr 3

Bit	7	6	5	4	3	2	1	0
byte 1	MQTT Control Packet type (3)				DUP flag	QoS level		RETAIN
	0	0	1	1	X	X	X	X
byte 2	Remaining Length							

- Topic Length
- The topic characters
- Packet ID

	Description	7	6	5	4	3	2	1	0
Topic Name									
byte 1	Length MSB (0)	0	0	0	0	0	0	0	0
byte 2	Length LSB (3)	0	0	0	0	0	0	1	1
byte 3	'a' (0x61)	0	1	1	0	0	0	0	1
byte 4	'/' (0x2F)	0	0	1	0	1	1	1	1
byte 5	'b' (0x62)	0	1	1	0	0	0	1	0
Packet Identifier									
byte 6	Packet Identifier MSB (0)	0	0	0	0	0	0	0	0
byte 7	Packet Identifier LSB (10)	0	0	0	0	1	0	1	0




Publish Message

- The QoS level specifies the answer

QoS Level	Expected Response
QoS 0	None
QoS 1	PUBACK Packet
QoS 2	PUBREC Packet

- I leave these for self study:
 - PUBACK – Publish acknowledgement (QoS 1)
 - PUBREC – Publish received (QoS 2 publish received, part 1)
 - PUBREL – Publish release (QoS 2 publish received, part 2)
 - PUBCOMP – Publish complete (QoS 2 publish received, part 3)



Subscribe Message

- Message nr 8


Bit	7	6	5	4	3	2	1	0
byte 1	MQTT Control Packet type (8)				Reserved			
	1	0	0	0	0	0	1	0
byte 2	Remaining Length							

- Package ID

	Description	7	6	5	4	3	2	1	0
Packet Identifier									
byte 1	Packet Identifier MSB (0)	0	0	0	0	0	0	0	0
byte 2	Packet Identifier LSB (10)	0	0	0	0	1	0	1	0

- Payload
 - Can have many subs in one message

Description	7	6	5	4	3	2	1	0
Topic Filter								
byte 1	Length MSB							
byte 2	Length LSB							
bytes 3..N	Topic Filter							
Requested QoS								
	Reserved						QoS	
byte N+1	0	0	0	0	0	0	X	X



Subscribe Ack Message

- Message nr 9

Bit	7	6	5	4	3	2	1	0
byte 1	MQTT Control Packet type (9)				Reserved			
	1	0	0	1	0	0	0	0
byte 2	Remaining Length							

- Packet ID to ACK


Bit	7	6	5	4	3	2	1	0
byte 1	Packet Identifier MSB							
byte 2	Packet Identifier LSB							

- Return code

Bit	7	6	5	4	3	2	1	0
	Return Code							
byte 1	X	0	0	0	0	0	X	X

Allowed return codes:

0x00 - Success - Maximum QoS 0
 0x01 - Success - Maximum QoS 1
 0x02 - Success - Maximum QoS 2
 0x80 - Failure



Unsubscribe Message

- Message nr 10


Bit	7	6	5	4	3	2	1	0
byte 1	MQTT Control Packet type (10)				Reserved			
	1	0	1	0	0	0	1	0
byte 2	Remaining Length							

- Packet ID

Bit	7	6	5	4	3	2	1	0
byte 1	Packet Identifier MSB							
byte 2	Packet Identifier LSB							

- Payload
 - Can have many unsubs in one message

	Description	7	6	5	4	3	2	1	0
Topic Filter									
byte 1	Length MSB (0)	0	0	0	0	0	0	0	0
byte 2	Length LSB (3)	0	0	0	0	0	0	1	1
byte 3	'a' (0x61)	0	1	1	0	0	0	0	1
byte 4	'/' (0x2F)	0	0	1	0	1	1	1	1
byte 5	'b' (0x62)	0	1	1	0	0	0	1	0




Unsubscribe Ack Message

- Message nr 11

Bit	7	6	5	4	3	2	1	0
byte 1	MQTT Control Packet type (11)				Reserved			
	1	0	1	1	0	0	0	0
byte 2	Remaining Length (2)							
	0	0	0	0	0	0	1	0

- Packet ID to ACK

Bit	7	6	5	4	3	2	1	0
byte 1	Packet Identifier MSB							
byte 2	Packet Identifier LSB							




PingReq/PingResp Message

- Ping Request

Bit	7	6	5	4	3	2	1	0
byte 1	MQTT Control Packet type (12)				Reserved			
	1	1	0	0	0	0	0	0
byte 2	Remaining Length (0)							
	0	0	0	0	0	0	0	0

- Ping Response

Bit	7	6	5	4	3	2	1	0
byte 1	MQTT Control Packet type (13)				Reserved			
	1	1	0	1	0	0	0	0
byte 2	Remaining Length (0)							
	0	0	0	0	0	0	0	0



Disconnect Message

- Disconnect

Bit	7	6	5	4	3	2	1	0
byte 1	MQTT Control Packet type (14)				Reserved			
	1	1	1	0	0	0	0	0
byte 2	Remaining Length (0)							
	0	0	0	0	0	0	0	0

V3.1.1 and V.5 and MQTT-SN

- MQTT version 3.1.1 was the standard for a long time
 - But has now been updated to v.5 (published 2019)
- In general it functions the same, but some changes to the protocol
 - Clean session/start
 - Client Restrictions/Limitations
 - Server Restrictions/Limitations
 - Will Delay Intervals
 - Server Redirect
 - Payload Format Indicator
 - Topic aliases
 - User Properties
 - Request Response
 - Non local publishing
 - Retained Message Control
 - Subscription Identifier
 - Shared Subscriptions
 - Reason Codes on All ACK Messages
 - Server Disconnect
- MQTT for Sensor Networks (MQTT-SN)
 - A more limited version of MQTT over UDP (not updated since 2013)



Mittuniversitetet
MID SWEDEN UNIVERSITY

Tips for the Lab Project

Tips for the Lab Project

- Start by using a normal MQTT Client, connect to an MQTT broker
 - For example `broker.mqttdashboard.com`
 - Observe the messages in Wireshark (`tcp.port == 1883`)
- Then create your own TCP socket listening on port 1883
 - And send your MQTT client to it instead
 - Observe, listening, and answer to incoming packets
- Since you are the broker, most answers will be short ACKs
 - The Connection ACK is only 4 bytes, for example:
 - Byte 1: (0010 0000) Packet Type 2, not DUP, QoS 0, no RETAIN
 - Byte 2: (0000 0010) 2 remaining bytes
 - Byte 3: (0000 0000) No session present
 - Byte 4: (0000 0000) return code 0, success

Tips for the Lab Project

- Start by answering the connect packages
 - And then make your program answer the ping packets
 - Otherwise all your clients will time out after a while
- If you receive a subscribe, save that socket and topic to a list/map
 - When you receive a publish, send it to all sockets on the topic list
 - Unsubscribe removes socket from the topic list
- Wireshark is your friend
 - To see how other MQTT clients messages looks like is very good for you to learn, debug and compare to yourself to them
- There is no delimiter of payload or between messages
 - Make sure you have the right remainder packet length
 - It is the only way for the system to distinguish between two messages

Tips for the Lab Project

- The MQTT 3.1.1 Oasis Standard
 - <http://docs.oasis-open.org/mqtt/mqtt/v3.1.1/mqtt-v3.1.1.html>
- A good online reference:
 - <http://www.steves-internet-guide.com/mqtt-basics-course/>

Contact Information

STEFAN FORSSTRÖM

PhD in Computer and System Sciences

MID SWEDEN UNIVERSITY

Institution of Information Systems and Technology (IST)

Campus Sundsvall, Room L426

Email: stefan.forsstrom@miun.se